


# Lecture 7

Binary Search Trees and Red-Black Trees

# Announcements

- HW 3 released! (Due Friday)


# Today: binary search trees



- Brief foray into **data structures!**
  - See CS 166 for more!
- What are binary search trees?
  - You may remember these from CS 106B
  - Why are they good?
  - Why are they bad?

this will lead us to...

- Self-Balancing Binary Search Trees
  - **Red-Black** trees.




# Why are we studying self-balancing BSTs?

1. The punchline is **important**:
  - A data structure with  $O(\log(n))$   
INSERT/DELETE/SEARCH
2. The idea behind **Red-Black Trees** is clever
  - It's good to be exposed to clever ideas.
  - Also it's just aesthetically pleasing.

# Motivation for binary search trees


- We've been assuming that we have access to some basic data structures:
  - (Sorted) linked lists




- (Sorted) arrays




# Sorted linked lists



- $O(1)$  insert/delete (assuming we have a pointer to the location of the insert/delete):



- $O(n)$  search/select:



# Sorted Arrays

1	2	3	4	5	7	8
---	---	---	---	---	---	---

- $O(n)$  insert/delete:

1	2	3	4	4.5	7	8
---	---	---	---	-----	---	---

- $O(\log(n))$  search,  $O(1)$  select:







1	2	3	4	5	7	8
---	---	---	---	---	---	---

Search: Binary search to see if 3 is in A.

Select: Third smallest is A[3].

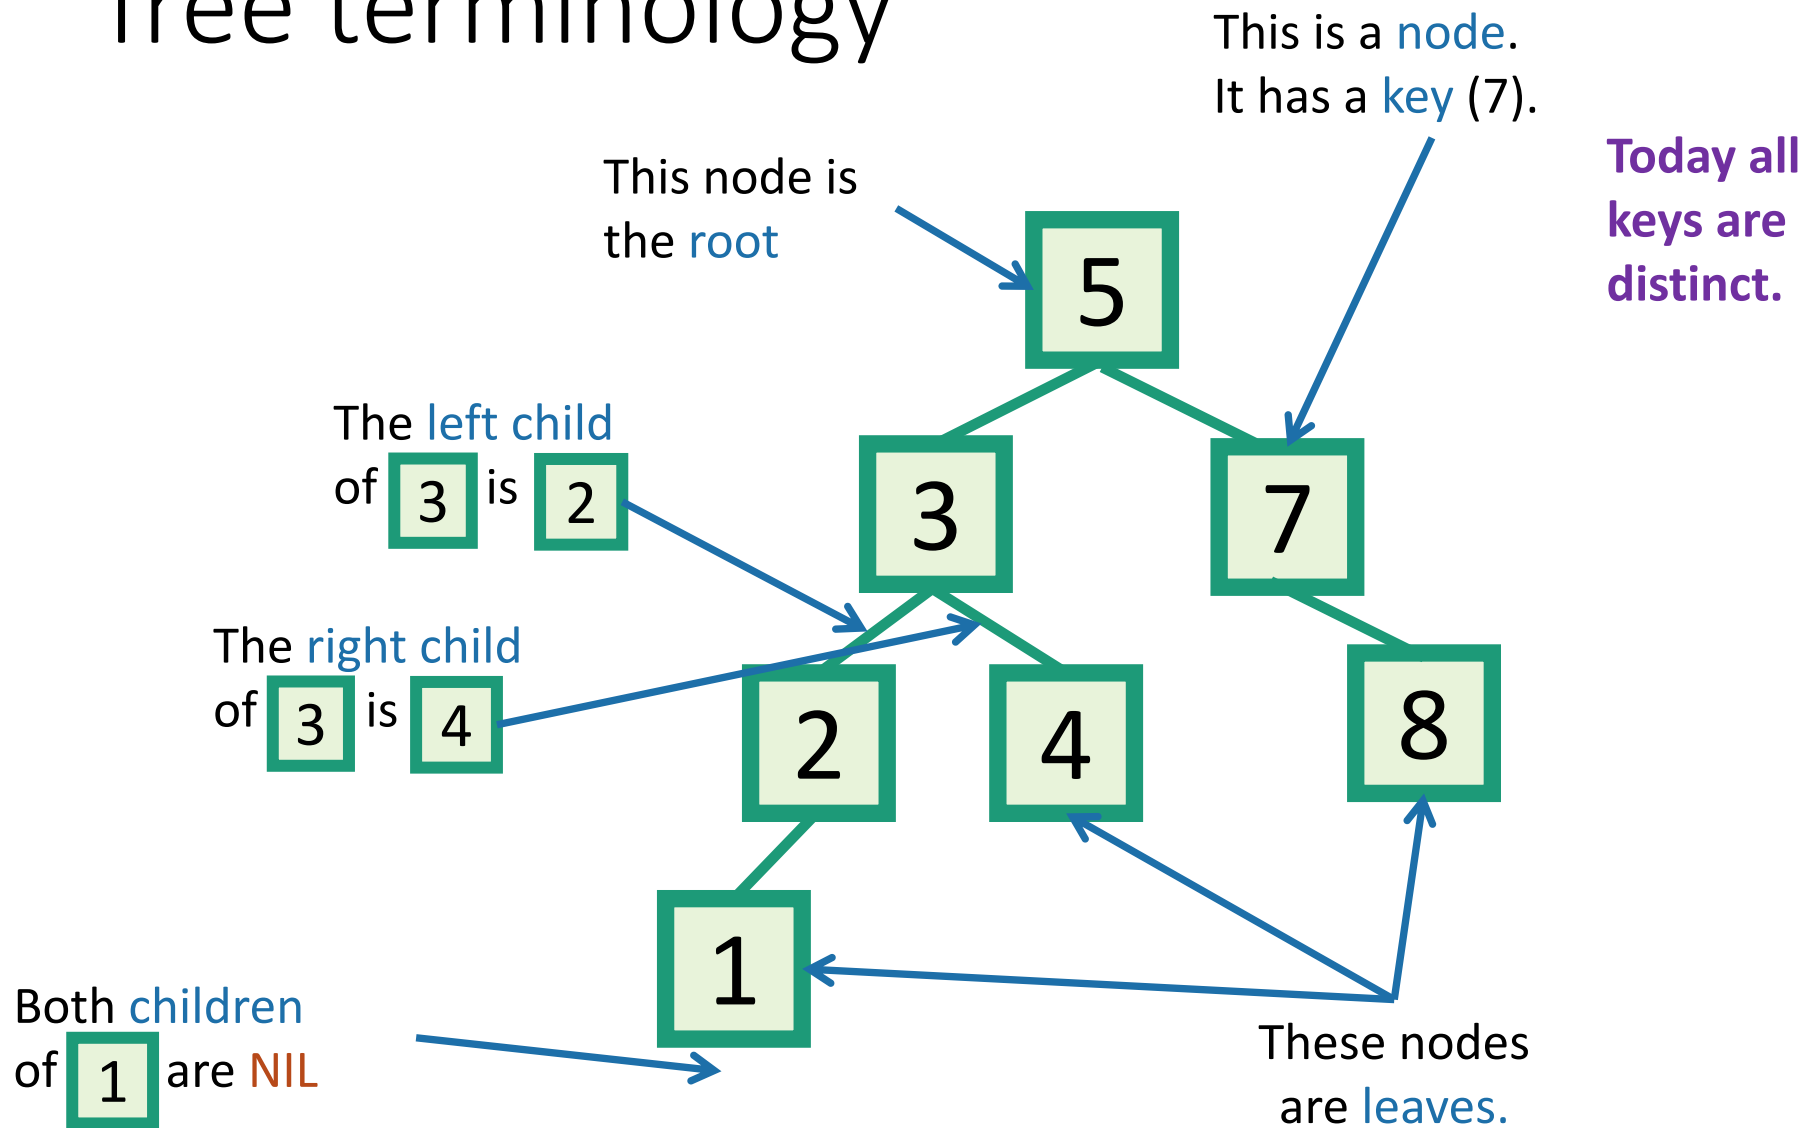
# The best of both worlds

**TODAY!**

	Sorted Arrays	Linked Lists	Binary Search Trees*
Search	$O(\log(n))$ 	$O(n)$ 	$O(\log(n))$ 
Insert/Delete	$O(n)$ 	$O(1)$ 	$O(\log(n))$ 




# Tree terminology




# Binary Search Trees

- It's a **binary tree** so that:
  - Every LEFT descendent of a node has key less than that node.
  - Every RIGHT descendent of a node has key larger than that node.
- Example of building a binary search tree:




# Binary Search Trees

- It's a **binary tree** so that:
  - Every LEFT descendent of a node has key less than that node.
  - Every RIGHT descendent of a node has key larger than that node.
- Example of building a binary search tree:




# Binary Search Trees

- It's a **binary tree** so that:
  - Every LEFT descendent of a node has key less than that node.
  - Every RIGHT descendent of a node has key larger than that node.
- Example of building a binary search tree:




# Binary Search Trees

- It's a **binary tree** so that:
  - Every LEFT descendent of a node has key less than that node.
  - Every RIGHT descendent of a node has key larger than that node.
- Example of building a binary search tree:




# Binary Search Trees

- It's a **binary tree** so that:
  - Every LEFT descendent of a node has key less than that node.
  - Every RIGHT descendent of a node has key larger than that node.
- Example of building a binary search tree:





**Aside:** this should look familiar  
kinda like QuickSort



# Binary Search Trees


- It's a **binary tree** so that:
  - **Every LEFT descendent** of a node has key less than that node.
  - **Every RIGHT descendent** of a node has key larger than that node.





# SEARCH in a Binary Search Tree

definition by example




**EXAMPLE:** Search for 4.

**EXAMPLE:** Search for 4.5


- It turns out it will be convenient to **return 4** in this case
- (that is, **return** the last node before we went off the tree)

Write pseudocode  
(or actual code) to  
implement this!



Ollie the over-achieving ostrich


# INSERT in a Binary Search Tree



## EXAMPLE: Insert 4.5

- **INSERT**(key):
  - $x = \text{SEARCH}(\text{key})$
  - **if**  $\text{key} > x.\text{key}$ :
    - Make a new node with the correct key, and put it as the right child of  $x$ .
  - **if**  $\text{key} < x.\text{key}$ :
    - Make a new node with the correct key, and put it as the left child of  $x$ .
  - **if**  $x.\text{key} == \text{key}$ :
    - **return**

# DELETE in a Binary Search Tree



## EXAMPLE: Delete 2


- **DELETE**(key):
  - $x = \text{SEARCH}(\text{key})$
  - **if**  $x.\text{key} == \text{key}$ :
    - ....delete  $x$ ....

This is a bit more complicated...


# DELETE in a Binary Search Tree

several cases (by example)


say we want to delete 3



**Case 1:** if 3 is a leaf,  
just delete it.



**Case 2:** if 3 has just one child,  
move that up.




We won't write  
down pseudocode  
for this – try to do it  
yourself!

Ollie the over-achieving ostrich

# DELETE in a Binary Search Tree ctd.


**Case 3:** if 3 has two children,  
replace 3 with its **immediate successor**.  
(aka, next biggest thing after 3)




- How do we find the immediate successor?
  - `SEARCH(3.right, 3)`
- How do we remove it when we find it?
  - Run DELETE for one of the previous two cases.
- Wait, what if it's **THIS** case? (Case 3).
  - It's not.

# How long do these operations take?

- **SEARCH** is the big one.
- Everything else just calls **SEARCH** and then does some small  $O(1)$ -time operation.



# Wait...



- This is a valid binary search tree.
- The version with  $n$  nodes has depth  $n$ , **not**  $O(\log(n))$ .

**Could such a tree show up?**  
In what order would I have to  
insert the nodes?

Inserting in the order  
2,3,4,5,6,7,8 would do it.

So this could happen.

# What to do?

How often is “every so often” in the worst case? It’s actually pretty often!



Ollie the over-achieving ostrich




- Keep track of how deep the tree is getting.
- If it gets too tall, re-do everything from scratch.
  - At least  $\Omega(n)$  every so often....
- Other ideas?





# Idea 1: Rotations

No matter what lives underneath A,B,C,  
this takes time  $O(1)$ . (Why?)

- Maintain Binary Search Tree (BST) property, while moving stuff around.



This seems helpful



# Does this work?

- Whenever something seems unbalanced, do rotations until it's okay again.



Even for me this is pretty vague. What do we mean by “seems unbalanced”? What’s “okay”?

Lucky the Lackadaisical Lemur

# Idea 2: have some proxy for balance

- Maintaining **perfect balance** is too hard.
- Instead, come up with some **proxy for balance**:
  - If the tree satisfies **[SOME PROPERTY]**, then it's pretty balanced.
  - We can maintain **[SOME PROPERTY]** using rotations.



There are actually several ways to do this, but today we'll see...


# Red-Black Trees

- A Binary Search Tree that balances itself!
- No more time-consuming by-hand balancing!
- Be the envy of your friends and neighbors with the time-saving...

*Red-Black tree!*

Maintain balance by stipulating that **black nodes** are balanced, and that there aren't too many **red nodes**.


*It's just good sense!*



# Red-Black Trees





these rules are the proxy for balance

- Every node is colored **red** or **black**.
- The root node is a **black node**.
- NIL children count as **black nodes**.
- Children of a **red node** are **black nodes**.
- For all nodes x:
  - all paths from x to NONE's have the same number of black nodes on them.




# Examples(?)

- Every node is colored **red** or **black**.
- The root node is a **black node**.
- NONE children count as **black nodes**.
- Children of a **red node** are **black nodes**.
- For all nodes x:
  - all paths from x to NONE's have the same number of black nodes on them.



# Why??????

- This is pretty balanced.
  - The **black nodes** are balanced
  - The **red nodes** are “spread out” so they don’t mess things up too much.
- We can maintain this property as we insert/delete nodes, by using rotations.





A cartoon illustration of a raccoon standing on its hind legs, smiling, with its arms outstretched. The raccoon has a grey body, a white chest, and a striped tail. It is positioned in the center of the page, above the text 'Raccoon'.

A cartoon illustration of a raccoon standing on its hind legs, smiling, with its arms outstretched. The raccoon has a grey body, a white chest, and a striped tail. It is positioned in the center of the page, above the 'Introduction' section.

A cartoon illustration of a raccoon standing on its hind legs, smiling, with its arms outstretched. The raccoon has a grey body, a white chest, and a striped tail. It is positioned in the center of the page, above the 'Introduction' section.

- 
- A cartoon illustration of a raccoon standing on its hind legs, smiling, with its arms outstretched. The raccoon has a grey body, a white chest, and a striped tail. It is positioned in the center of the page, above the 'Introduction' section.



A cartoon illustration of a raccoon standing on its hind legs, smiling, with its arms outstretched. The raccoon has a grey body, a white chest, and a striped tail. It is positioned in the center of the page, above the 'Introduction' section.

A cartoon illustration of a raccoon standing on its hind legs, smiling, with its arms outstretched. The raccoon has a grey body, a white chest, and a striped tail. It is positioned in the center of the page, above the 'Introduction' section.




# That turns out to be basically right.

[proof sketch]



- Say there are  $b(x)$  black nodes in **any path from  $x$  to NONE**.
  - (including  $x$ ).
- **Claim:**
  - Then there are at least  $2^{b(x)} - 1$  nodes in the subtree underneath  $x$ .
- [Proof by induction — on board if time]



Then:

$$n \geq 2^{b(\text{root})} - 1 \quad \text{using the Claim}$$

$$\geq 2^{\text{height}/2} - 1 \quad b(\text{root}) \geq \text{height}/2 \text{ because of RBTree rules.}$$

Rearranging:

$$n + 1 \geq 2^{\text{height}/2} \Rightarrow \text{height} \leq 2\log(n + 1)$$

Okay, so it's balanced...


...but can we maintain it?

- Yes!


- For the rest of lecture:
  - sketch of how we'd do this.
- See CLRS for more details.

# Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.



What if it looks like this?




Example: insert 0




# Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.




What if it looks like this?

Example: insert 0




# Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.
- **Fix things up if needed.**



What if it looks like this?



Example: insert 0



Can't we just insert 0 as a **black node**?



# Inserting into a Red-Black Tree

-1

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.




What if it looks like this?

Example: insert 0


- Instead recolor like this.
- Need to argue:
  - RB-Tree properties still hold.
- What about the red root?
  - if 6 is actually the root, color it black.
  - Else, recursively re-color up the tree.



Now the problem looks like this, where I'm inserting




# Inserting into a Red-Black Tree



What if it looks like this?

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.




Example: Insert 0.

- Actually, **this can't happen?**
- It might happen that we just turned 0 red from the previous step.
- Or it could happen if **7** is actually NIL.




# Recall Rotations

- Maintain Binary Search Tree (BST) property, while moving stuff around.





# Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.



What if it looks like this?

**YOINK!**



Need to argue that if RB-Tree property held before, it still does.

# That's basically it

- A few things still left to check for **INSERT!**
  - Anything else that might happen looks basically like what we just did.
  - Formally dealing with the recursion.
  - **You check these!** (or see CLRS)
- **DELETE** is similar.









Plucky the pedantic penguin

## The punchline:

- Red-Black Trees **always** have height at most  $2\log(n+1)$ .
- As with general **Binary Search Trees**, all operations are  $O(\text{height})$
- **So all operations are  $O(\log(n))$ .**

# Conclusion: The best of both worlds

	Sorted Arrays	Linked Lists	Balanced Binary Search Trees
Search	$O(\log(n))$ 	$O(n)$ 	$O(\log(n))$ 
Insert/Delete	$O(n)$ 	$O(1)$ 	$O(\log(n))$ 

# Recap

- **Balanced binary trees** are the best of both worlds!
- But we need to **keep them balanced**.
- **Red-Black Trees** do that for us.
  - We get  $O(\log(n))$ -time INSERT/DELETE/SEARCH
  - Clever idea: have a **proxy for balance**

## Next time

- **Hashing!**